


Functional requirements

- Match buyers & sellers of stock!
- ordered set of orders that everybody agrees on
 - eventually!
 - fault tolerant

diff from stock brokers!

General algorithm



The limit order book

(matching algo)

can collect orders placed at etc.

Keeping track of bids & asks

Bids

Asks

BK: \$100 x 5
MB: \$90 x 10
C: 120 \$ x 100
M: 130 \$ x 10
R: 140 \$ x 20

one or more trades
exec when

bid price $\geq$ ask price

X Trades

✓ 130 x 110 | bid
120 x 100 |
10 x 10 | sell

Performance is key: $\sim 10^6$ orders daily.

Low throughputs result in many of these trades

to start crawling

{ Matching bids to asks should be at high throughputs }

Limit I/O as much as possible.

(n/w & disk ops)

Limit loads on I/O.

Deterministic req: $S \xrightarrow{\{\text{set of } m\} \text{ ops}} S'$
1 $\leftrightarrow$ 1 connection, threads run on same CPos and h/w
 $\uparrow$, can't mix to all at once

TCP (X)

UDP multicast (✓)

- msg drop possible

- msg rcv out of order

Multicast networking

or communication is essential here. $\therefore$ same data needs to be sent to all clients.


I/O is to be avoided hence

commit logs shouldn't be used.




or can't. Use multicast to dump msgs in commit log (multiple) and clients can read them whenever they want?

IGN + Layer 2 Collects data from stream
AM Layer 3




connection should be "of air".

Data should reach at any of the clients at same time w/ minimal latency.



Use "retransmission"

m_0, m_1, m_2, m_3



- caches msgs

- a lot of RTs used


to reduce P (dropping messages)

- RTs communicate internally as well to sync

fault tolerance via state m/c replicatⁿ

- runs in n/m. Backups are essential.
everything


Initial state copied to P




Whenever ① goes down ② takes over. And new
 (R), (L2) pulls last state from (RTSx) and
 replays msgs after it via (RTx) &
SAFA algo: to develop fresh state.

UDP broad (R)
 cast

Zookeeper with heartbeats to
 resolve problem.



order is possibly diff on (P) & (R).



multicast should be trigger for all transact^{ns}
 even if we're using ① for Ack. and
 telling everyone about offered trades.

For ① we'll need $RT(S)$ & $RT(S_i)$.
②

Partitioning: we're relying on a single leader,
 pa. . . is essential.

Partitioning on multiple stock collections on a
 single server is essential. ordering can't be
 maintained otherwise.

↳ Race condⁿ and order violations possible. ↳ I/O would be introduced
 DB interact^{ns} n/w

128/256 GB m/c's are needed.

Atomicity on multiple symbols need
 a slow distributed transactions.

ordering on matching engine

Multi-threaded env is still there.


of ip $\Rightarrow$ grab lock on relevant client/DS

[exec BID AAPL 1000 20 client 1
ASK AAPL 1020 30 C2
BID AAPL 1010 10 C3] id=1001

Lock on C1, C2, C3 reqd. to get a
reqd. seq number.

[cancel ASK AAPL 1020 30 C2] id=1002
will be exc ~~after~~ after 1001
or before

not simultaneously.



YOUTUBE SYS DESIGN

- upload + download/watcher ★
- Search ★
- Recommend [a separate topic]
- Commenting ["]
- Analytics [a separate topic]

Non funct. req.

Scale : 1B users
400 M DAU

Streams
Downloads
avg : 10^4 : 1 uploads
(assumed)


week

no. of uploads per day : $10^9 \times 0.2$

$$\Rightarrow 1.26B \times 10^9 \times 0.1 \times 60$$

$\Rightarrow$

Functionality



1 video = 3GB

↓
1.2GB